

Kubernetes on Joyent

The Hard Way

The bulk of this article is Copyrighted by Kelsey Hightower and used under the Apache license. The original text and license can be found on [Github](#)

Kelsey's [Kubernetes the Hard Way](#) is regarded in the Kubernetes community as one of the best ways to learn how to deploy Kubernetes by hand, prior to engaging in an automation strategy. Since this work is so well known, we've used it as our template for this guide.

Cloud Infrastructure Provisioning

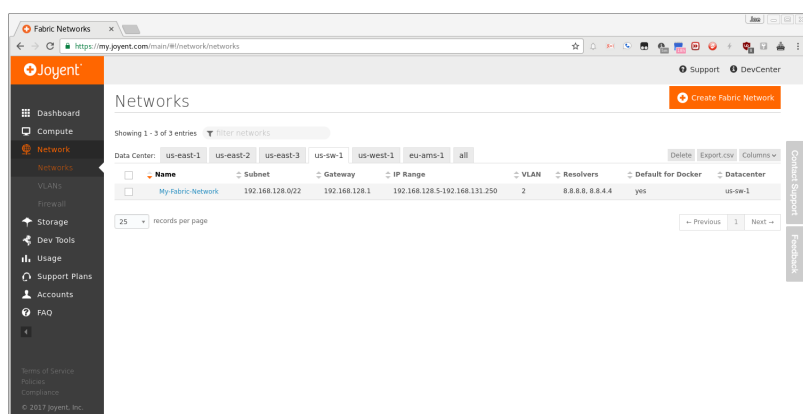
This tutorial shows how to run an example HA Kubernetes cluster using the Joyent public cloud.

Begin by creating a [My Joyent](#) account, and [installing and configuring the Triton CLI](#).

This tutorial assumes the us-sw-1 data center is being used.

Network

Create a private Joyent Fabric Network for your Kubernetes installation.



From the [networks](#) screen, select [Create Fabric Network](#)

Create Fabric Network

Name:

Data Center:

Subnet:
Subnet needs to be specified in CIDR Notation.

Gateway:
Gateway IP address

IP Range:

VLAN:
[Create a new VLAN](#)

DNS Resolvers:
Resolver IP addresses, use Enter button as delimiter

Routes:
[Add another route](#)

Description:

☒ Provision a NAT zone on the gateway address.

* required fields

[Cancel](#) [Create Network](#)

Create a network with these attributes:

- Name: kubernetes
- Data Center: us-sw-1
- Subnet: 10.240.0.0/24
- Gateway: 10.240.0.1
- IP Range: 10.240.0.2 10.240.0.254
- VLAN: *Leave at default*
- DNS Resolvers: 8.8.8.8 8.8.4.4
- Provision a NAT zone on the gateway address: *Checked*

Provision Virtual Machines

Provision the compute instances required for running a HA Kubernetes cluster. All the VMs in this tutorial are provisioned using Ubuntu 16.04 because it runs a recent Linux Kernel that has good support for Docker. A total of seven virtual machines will be created.

Virtual Machines

Bastion

```
triton instance create \
  --wait \
  --name=kubernetes \
  -N Joyent-SDC-Public,kubernetes \
  -m user-data="hostname=kubernetes" \
  ubuntu-16.04 g4-highcpu-16
```

Controllers

```
triton instance create \
  --wait \
  --name=controller0 \
  -N kubernetes \
  -m user-data="hostname=controller0" \
  ubuntu-16.04 g4-highcpu-16
```

```
triton instance create \
  --wait \
  --name=controller1 \
  -N kubernetes \
  -m user-data="hostname=controller1" \
  ubuntu-16.04 g4-highcpu-16
```

```
triton instance create \  
  --wait \  
  --name=controller2 \  
  -N kubernetes \  
  -m user-data="hostname=controller2" \  
  ubuntu-16.04 g4-highcpu-16
```

Worker Agents

```
triton instance create \  
  --wait \  
  --name=worker0 \  
  -N kubernetes \  
  -m user-data="hostname=worker0" \  
  ubuntu-certified-16.04 k4-highcpu-kvm-1.75G
```

```
triton instance create \  
  --wait \  
  --name=worker1 \  
  -N kubernetes \  
  -m user-data="hostname=worker0" \  
  ubuntu-certified-16.04 k4-highcpu-kvm-1.75G
```

```
triton instance create \  
  --wait \  
  --name=worker2 \  
  -N kubernetes \  
  -m user-data="hostname=worker0" \  
  ubuntu-certified-16.04 k4-highcpu-kvm-1.75G
```

Allow root login to certified image:

```
BASTION=$(triton ip kubernetes)  
IP=$(triton instance list -o 'ips' -j name=worker0 | sed 's/.*ips":.'"\[0-9.\]\+\)".*/\1/')  
ssh -o proxycommand="ssh root@$BASTION -W %h:%p" ubuntu@$IP cp  
/home/ubuntu/.ssh/authorized_keys /root/.ssh/  
IP=$(triton instance list -o 'ips' -j name=worker1 | sed 's/.*ips":.'"\[0-9.\]\+\)".*/\1/')  
ssh -o proxycommand="ssh root@$BASTION -W %h:%p" ubuntu@$IP cp  
/home/ubuntu/.ssh/authorized_keys /root/.ssh/  
IP=$(triton instance list -o 'ips' -j name=worker2 | sed 's/.*ips":.'"\[0-9.\]\+\)".*/\1/')  
ssh -o proxycommand="ssh root@$BASTION -W %h:%p" ubuntu@$IP cp  
/home/ubuntu/.ssh/authorized_keys /root/.ssh/
```

You should now have the following compute instances:

```
triton ls
```

SHORTID	NAME	IMG	STATE	FLAGS	AGE
e2a3b8a6	kubernetes	ubuntu-16.04@20161213	running	-	1d
b080b2c4	controller0	ubuntu-16.04@20161213	running	-	1d
846d3bb7	controller1	ubuntu-16.04@20161213	running	-	1d
e3226afb	controller2	ubuntu-16.04@20161213	running	-	1d
753c5f0c	worker0	ubuntu-certified-16.04@20161221	running	K	1d
7bb2a272	worker1	ubuntu-certified-16.04@20161221	running	K	1d
c34597cf	worker2	ubuntu-certified-16.04@20161221	running	K	1d

Setting up a Certificate Authority and TLS Cert Generation

Set up the necessary PKI infrastructure to secure the Kubernetes components. This tutorial uses CloudFlare's PKI toolkit, [cfssl](#), to bootstrap a Certificate Authority and generate TLS certificates.

You will generate a single set of TLS certificates that can be used to secure the following Kubernetes components:

- etcd
- Kubernetes API Server
- Kubernetes Kubelet

Best Practice

NOTE In production you should strongly consider generating individual TLS certificates for each component.

After completing this section you should have the following TLS keys and certificates:

```
ca-key.pem
ca.pem
kubernetes-key.pem
kubernetes.pem
```

Install CFSSL

This tutorial requires the [cfssl](#) and [cfssljson](#) binaries. Download them from the [cfssl repository](#).

OS X

```
wget -O /usr/local/bin/cfssl https://pkg.cfssl.org/R1.2/cfssl_darwin-amd64
wget -O /usr/local/bin/cfssljson https://pkg.cfssl.org/R1.2/cfssljson_darwin-amd64
chmod +x /usr/local/bin/cfssl
chmod +x /usr/local/bin/cfssljson
```

Linux

```
wget -O /usr/local/bin/cfssl https://pkg.cfssl.org/R1.2/cfssl_linux-amd64
wget -O /usr/local/bin/cfssljson https://pkg.cfssl.org/R1.2/cfssljson_linux-amd64
chmod +x /usr/local/bin/cfssl
chmod +x /usr/local/bin/cfssljson
```

Set up a Certificate Authority

Create the CA configuration file:

```
echo '{
  "signing": {
    "default": {
      "expiry": "8760h"
    },
    "profiles": {
      "kubernetes": {
        "usages": ["signing", "key encipherment", "server auth", "client auth"],
        "expiry": "8760h"
      }
    }
  }
}' > ca-config.json
```

Generate the CA certificate and private key

Create the CA CSR:

```
echo '{
  "CN": "Kubernetes",
  "key": {
    "algo": "rsa",
    "size": 2048
  },
  "names": [
    {
      "C": "US",
      "L": "Portland",
      "O": "Kubernetes",
      "OU": "CA",
      "ST": "Oregon"
    }
  ]
}' > ca-csr.json
```

Generate the CA certificate and private key:

```
cfssl gencert -initca ca-csr.json | cfssljson -bare ca
```

Results

```
ca-key.pem
ca.csr
ca.pem
```

Verification

```
openssl x509 -in ca.pem -text -noout
```

Generate the single Kubernetes TLS Cert

Generate a TLS certificate that will be valid for all Kubernetes components. Using a single certificate will simplify this tutorial. However, in production you should generate individual TLS certificates for each component type. All replicas of a given component type must share the same certificate.

Set the Kubernetes Public Address:

```
IPS=$(triton instance ls -o ips | grep -v IPS | tr '\n' ',' | tr -d '[]')
```

Create the `kubernetes-csr.json` file:

```
cat > kubernetes-csr.json <<EOF
{
  "CN": "kubernetes",
  "hosts": [
    "kubernetes",
    "${IPS}"
    "127.0.0.1"
  ],
  "key": {
    "algo": "rsa",
    "size": 2048
  },
  "names": [
    {
      "C": "US",
      "L": "Portland",
      "O": "Kubernetes",
      "OU": "Cluster",
      "ST": "Oregon"
    }
  ]
}
EOF
```

Generate the Kubernetes certificate and private key:

```
cfssl gencert \
  -ca=ca.pem \
  -ca-key=ca-key.pem \
  -config=ca-config.json \
  -profile=kubernetes \
  kubernetes-csr.json | cfssljson -bare kubernetes
```

Results

```
kubernetes-key.pem
kubernetes.csr
kubernetes.pem
```

Verification

```
openssl x509 -in kubernetes.pem -text -noout
```


Copy TLS Certs

Copy the TLS certificates and keys to each Kubernetes host:

```
KUBERNETES_HOSTS=(controller0 controller1 controller2 worker0 worker1 worker2)
BASTION=$(triton ip kubernetes)
tar cf - ca.pem kubernetes-key.pem kubernetes.pem |
  ssh $BASTION tar xf -
for ip in $(echo $IPS | tr ',' ' '); do
  tar cf - ca.pem kubernetes-key.pem kubernetes.pem |
  ssh -o StrictHostKeyChecking=no \
  -o ProxyCommand="ssh root@$BASTION -W %h:%p"
  root@$ip tar xf -
done
```

Bootstrapping a HA etcd cluster

Bootstrap a three node etcd cluster with the following virtual machines:

- controller0
- controller1
- controller2

Considerations for Running etcd

All Kubernetes components are stateless, which greatly simplifies managing a Kubernetes cluster. All state is stored in etcd, which is a database and must be treated specially. To limit the number of compute resource needed in this tutorial, etcd is installed on the Kubernetes controller nodes. In production environments, etcd should be run on a dedicated set of machines for the following reasons:

- The etcd lifecycle is not tied to Kubernetes. You should be able to upgrade etcd independently of Kubernetes.
- You should be able to scale out etcd independently of Kubernetes.
- It prevents other applications from consuming resources (CPU, Memory, I/O) required by etcd.

Provision the etcd Cluster

The provisioning commands must be run on each of the etcd machines: `controller0`, `controller1`, and `controller2`.

TLS Certificates

The TLS certificates created in the [Setting up a CA and TLS Cert Generation](#) section are used to secure communication between the Kubernetes API server and the etcd cluster. The TLS certificates

also limit access to the etcd cluster using TLS client authentication. Only clients with a TLS certificate signed by a trusted CA are able to access the etcd cluster.

Copy the TLS certificates to the etcd configuration directory:

```
mkdir -p /etc/etcd/  
cp ca.pem kubernetes-key.pem kubernetes.pem /etc/etcd/
```

Download and Install the etcd binaries

Download the official etcd release binaries from [coreos/etcd](https://github.com/coreos/etcd) GitHub project:

```
wget https://github.com/coreos/etcd/releases/download/v3.0.10/etcd-v3.0.10-linux-  
amd64.tar.gz
```

*Extract and install the **etcd** server binary and the **etcdctl** command line client:*

```
tar -xvf etcd-v3.0.10-linux-amd64.tar.gz  
mv etcd-v3.0.10-linux-amd64/etcd\* /usr/bin/
```

All etcd data is stored under the etcd data directory.

NOTE

Best Practice

In production, the etcd data directory should be backed by a persistent disk.

Copy the TLS certificates to the etcd configuration directory:

```
mkdir -p /var/lib/etcd  
cp ca.pem kubernetes-key.pem kubernetes.pem /etc/etcd/
```

The internal IP address is used by etcd to serve client requests and communicate with other etcd peers.

Set The Internal IP Address:

```
INTERNAL_IP=$(ip addr show eth0 | awk '/inet /{gsub(/\./[0-9][0-9]/,"");print $2}')
```

Each etcd member must have a unique name within an etcd cluster

```
ETCD_NAME=$(mdata-get user-data | sed 's/.*hostname=([^\ ]*)\.*$/\1/')
```

Set the controller addresses:

```
INITIAL_CLUSTER=$(triton instance ls -o name,ips | awk '/controller/{gsub(/\./[0-9][0-9]/,""),$2);print $1 "=https://" $2":2380"}' | tr '\n' ',')
```

The etcd server is started and managed by systemd.

Create the etcd systemd service file:

```
echo "[Unit]
Description=etcd
Documentation=https://github.com/coreos

[Service]
ExecStart=/usr/bin/etcd --name $ETCD_NAME \
  --cert-file=/etc/etcd/kubernetes.pem \
  --key-file=/etc/etcd/kubernetes-key.pem \
  --peer-cert-file=/etc/etcd/kubernetes.pem \
  --peer-key-file=/etc/etcd/kubernetes-key.pem \
  --trusted-ca-file=/etc/etcd/ca.pem \
  --peer-trusted-ca-file=/etc/etcd/ca.pem \
  --initial-advertise-peer-urls https://$INTERNAL_IP:2380 \
  --listen-peer-urls https://$INTERNAL_IP:2380 \
  --listen-client-urls https://$INTERNAL_IP:2379,http://127.0.0.1:2379 \
  --advertise-client-urls https://$INTERNAL_IP:2379 \
  --initial-cluster-token etcd-cluster-0 \
  --initial-cluster $INITIAL_CLUSTER \
  --initial-cluster-state new \
  --data-dir=/var/lib/etcd
Restart=on-failure
RestartSec=5

[Install]
WantedBy=multi-user.target" > /etc/systemd/system/etcd.service
```

Enable and start etcd:

```
systemctl daemon-reload
systemctl enable etcd
systemctl start etcd
```

Verification

```
systemctl status etcd --no-pager
```

NOTE Remember to run these commands on `controller0`, `controller1`, and `controller2`

Verification

Once all three etcd nodes have been bootstrapped, verify the etcd cluster is healthy. Request etcd cluster health on one of the controller nodes:

```
etcdctl --ca-file=/etc/etcd/ca.pem cluster-health
```

```
member 3a57933972cb5131 is healthy: got healthy result from https://10.240.0.12:2379
member f98dc20bce6225a0 is healthy: got healthy result from https://10.240.0.10:2379
member ffed16798470cab5 is healthy: got healthy result from https://10.240.0.11:2379
cluster is healthy
```

Bootstrapping a HA Kubernetes Control Plane

Bootstrap a three node Kubernetes controller cluster with the following virtual machines:

- controller0
- controller1
- controller2

Considerations for Running the Control Plane

The control plane is made up of the following Kubernetes components:

- Kubernetes API Server
- Kubernetes Scheduler
- Kubernetes Controller Manager

Each component is run on the same machine for the following reasons:

- The Scheduler and Controller Manager are tightly coupled with the API Server
- Only one Scheduler and Controller Manager can be active at a given time, but it's OK to run multiple instances at the same time. Each component will elect a leader via the API Server.
- Running multiple instances of each component is required for HA
- Running each component on the same machine as the API Server simplifies configuration.

Provision the Kubernetes Controller Cluster

Run the provisioning commands on `controller0`, `controller1`, and `controller2`.

Download and install the Kubernetes controller binaries:

```
wget -P /usr/local/bin/ https://storage.googleapis.com/kubernetes-  
release/release/v1.5.2/bin/linux/amd64/{kube-apiserver,kube-controller-manager,kube-  
scheduler,kubectl}  
chmod +x /usr/local/bin/{kube-apiserver,kube-controller-manager,kube-  
scheduler,kubectl}
```

TLS Certificates

The TLS certificates created in the [Setting up a CA and TLS Cert Generation](#) section are used to secure communication between the Kubernetes API server and Kubernetes clients such as **kubectl** and the **kubelet** agent. The TLS certificates is also used to authenticate the Kubernetes API server to etcd via TLS client authentication.

Copy the TLS certificates to the Kubernetes configuration directory:

```
mkdir -p /var/lib/kubernetes  
cp ca.pem kubernetes-key.pem kubernetes.pem /var/lib/kubernetes/
```

Setup Authentication and Authorization

Authentication

[Token based authentication](#) is used to limit access to the Kubernetes API. The authentication token is used by the following components:

- kubelet (client)
- Kubernetes API Server (server)

The other components, mainly the **scheduler** and **controller manager**, access the Kubernetes API server locally over the insecure API port which does not require authentication. The insecure port is only enabled for local access.

Create a token file:

```
COMMON_TOKEN=$(head /dev/urandom | base32 | head -c 8)  
echo "${COMMON_TOKEN},admin,admin  
${COMMON_TOKEN},scheduler,scheduler  
${COMMON_TOKEN},kubelet,kubelet  
" > /var/lib/kubernetes/token.csv  
echo "Your common token is: ${COMMON_TOKEN}"
```

Make a note of your common token for later.

Authorization

Attribute-Based Access Control (ABAC) will be used to authorize access to the Kubernetes API. In this tutorial ABAC is set up using the Kubernetes policy file backend as documented in the [Kubernetes authorization guide](#).

Create the authorization policy file:

```
cat << EOF > /var/lib/kubernetes/authorization-policy.jsonl
{"apiVersion": "abac.authorization.kubernetes.io/v1beta1", "kind": "Policy", "spec":
{"user": "*", "nonResourcePath": "*", "readOnly": true}}
{"apiVersion": "abac.authorization.kubernetes.io/v1beta1", "kind": "Policy", "spec":
{"user": "admin", "namespace": "*", "resource": "*", "apiGroup": ""}}
{"apiVersion": "abac.authorization.kubernetes.io/v1beta1", "kind": "Policy", "spec":
{"user": "scheduler", "namespace": "*", "resource": "*", "apiGroup": ""}}
{"apiVersion": "abac.authorization.kubernetes.io/v1beta1", "kind": "Policy", "spec":
{"user": "kubelet", "namespace": "*", "resource": "*", "apiGroup": ""}}
{"apiVersion": "abac.authorization.kubernetes.io/v1beta1", "kind": "Policy", "spec":
{"group": "system:serviceaccounts", "namespace": "*", "resource": "*", "apiGroup": "",
"nonResourcePath": ""}}
EOF
```

Kubernetes API Server

Create the systemd service file:

```
INTERNAL_IP=$(ip addr show eth0 | awk '/inet /{gsub(/\[/[0-9][0-9]/, "");print $2}')
ETCD_SERVERS=$(triton instance ls -o name,ips | awk '/controller/{gsub(/^[0-9.]/, "", $2);print "https://" $2":2380"}' | tr '\n' ',')
echo "[Unit]
Description=Kubernetes API Server
Documentation=https://github.com/GoogleCloudPlatform/kubernetes

[Service]
ExecStart=/usr/bin/kube-apiserver \
  --admission
-control=NamespaceLifecycle,LimitRanger,SecurityContextDeny,ServiceAccount,ResourceQuota \
  --advertise-address=${INTERNAL_IP} \
  --allow-privileged=true \
  --apiserver-count=3 \
  --authorization-mode=ABAC \
  --authorization-policy-file=/var/lib/kubernetes/authorization-policy.jsonl \
  --bind-address=0.0.0.0 \
  --enable-swagger-ui=true \
  --etcd-cafile=/var/lib/kubernetes/ca.pem \
  --insecure-bind-address=0.0.0.0 \
  --kubelet-certificate-authority=/var/lib/kubernetes/ca.pem \
  --etcd-servers=${ETCD_SERVERS} \
  --service-account-key-file=/var/lib/kubernetes/kubernetes-key.pem \
  --service-cluster-ip-range=10.32.0.0/24 \
  --service-node-port-range=30000-32767 \
  --tls-cert-file=/var/lib/kubernetes/kubernetes.pem \
  --tls-private-key-file=/var/lib/kubernetes/kubernetes-key.pem \
  --token-auth-file=/var/lib/kubernetes/token.csv \
  --v=2
Restart=on-failure
RestartSec=5

[Install]
WantedBy=multi-user.target" > /etc/systemd/system/kube-apiserver.service
```

Enable and start API server:

```
systemctl daemon-reload
systemctl enable kube-apiserver
systemctl start kube-apiserver
```

Verification

```
systemctl status kube-apiserver --no-pager
```

Kubernetes Controller Manager

```
echo "[Unit]
Description=Kubernetes Controller Manager
Documentation=https://github.com/GoogleCloudPlatform/kubernetes

[Service]
ExecStart=/usr/bin/kube-controller-manager \
  --allocate-node-cidrs=true \
  --cluster-cidr=10.200.0.0/16 \
  --cluster-name=kubernetes \
  --leader-elect=true \
  --master=http://$INTERNAL_IP:8080 \
  --root-ca-file=/var/lib/kubernetes/ca.pem \
  --service-account-private-key-file=/var/lib/kubernetes/kubernetes-key.pem \
  --service-cluster-ip-range=10.32.0.0/24 \
  --v=2
Restart=on-failure
RestartSec=5

[Install]
WantedBy=multi-user.target
" > /etc/systemd/system/kube-controller-manager.service
```

Enable and start the controller manager:

```
systemctl daemon-reload
systemctl enable kube-controller-manager
systemctl start kube-controller-manager
```

Verification

```
systemctl status kube-controller-manager --no-pager
```

Kubernetes Scheduler


```

echo "[Unit]
Description=Kubernetes Scheduler
Documentation=https://github.com/GoogleCloudPlatform/kubernetes

[Service]
ExecStart=/usr/bin/kube-scheduler \
  --leader-elect=true \
  --master=http://$INTERNAL_IP:8080 \
  --v=2
Restart=on-failure
RestartSec=5

[Install]
WantedBy=multi-user.target
" > /etc/systemd/system/kube-scheduler.service

```

Enable and start the scheduler:

```

systemctl daemon-reload
systemctl enable kube-scheduler
systemctl start kube-scheduler

```

Verification

```
systemctl status kube-scheduler --no-pager
```

Verification

```
kubectl get componentstatuses
```

NAME	STATUS	MESSAGE	ERROR
controller-manager	Healthy	ok	
scheduler	Healthy	ok	
etcd-1	Healthy	{"health": "true"}	
etcd-0	Healthy	{"health": "true"}	
etcd-2	Healthy	{"health": "true"}	

NOTE Remember to run these steps on **controller0**, **controller1**, and **controller2**

Setup Kubernetes API Server Frontend Load Balancer

Setting up a load balancer is out of the scope of this tutorial. A frontend load balancer, like haproxy, will need to be set up on bastion in order to access any of your APIs or services externally. If you don't want to set up haproxy yet, complete all the remaining tasks from the bastion.

Kubernetes Workers

Bootstrap three Kubernetes worker nodes on the following virtual machines:

- worker0
- worker1
- worker2

Considerations for Running Worker Nodes

Kubernetes worker nodes are responsible for running your containers. All Kubernetes clusters need one or more worker nodes. Worker nodes should be run on dedicated machines for the following reasons:

- Ease of deployment and configuration
- It avoids mixing arbitrary workloads with critical cluster components. You can build machines with just enough resources so you don't have to worry about waste.

Some people would like to run workers and cluster services anywhere in the cluster. This is totally possible, and you'll have to decide what's best for your environment.

Run the following commands on each of the three machines `worker0`, `worker1`, and `worker2`.

Move the TLS certificates in place

```
mkdir -p /var/lib/kubernetes  
cp ca.pem kubernetes-key.pem kubernetes.pem /var/lib/kubernetes/
```

Docker

Kubernetes should be compatible with the Docker 1.9.x - 1.12.x:

Download and install Docker:

```
wget https://get.docker.com/builds/Linux/x86_64/docker-1.12.1.tgz  
tar -xvf docker-1.12.1.tgz  
cp docker/docker* /usr/bin/
```

Create the Docker systemd service file:

```
sh -c 'echo "[Unit]
Description=Docker Application Container Engine
Documentation=http://docs.docker.io

[Service]
ExecStart=/usr/bin/docker daemon \
  --iptables=false \
  --ip-masq=false \
  --host=unix:///var/run/docker.sock \
  --log-level=error \
  --storage-driver=overlay
Restart=on-failure
RestartSec=5

[Install]
WantedBy=multi-user.target" > /etc/systemd/system/docker.service'
```

Enable and start Docker:

```
systemctl daemon-reload
systemctl enable docker
systemctl start docker
```

Verification

```
docker version
```

kubelet

Download and install CNI plugins:

```
mkdir -p /opt/cni
wget https://storage.googleapis.com/kubernetes-release/network-plugins/cni-
07a8a28637e97b22eb8dfe710eeae1344f69d16e.tar.gz
tar -xvf cni-07a8a28637e97b22eb8dfe710eeae1344f69d16e.tar.gz -C /opt/cni
```

Download and install the Kubernetes worker binaries:

```
wget -P /usr/bin https://storage.googleapis.com/kubernetes-
release/release/v1.5.2/bin/linux/amd64/{kubect1,kube-proxy,kubelet}
chmod +x /usr/bin/{kubect1,kube-proxy,kubelet}
```

Set **COMMON_TOKEN** to the value you saved in [Authentication](#).

Configure kubelet

```
CONTROLLER0=$(triton instance ls -o name,ips | awk '/controller0/{gsub(/^[^0-9.]/,"",$2);print "https://" $2 ":6443"}')
API_SERVERS=$(triton instance ls -o name,ips | awk '/controller/{gsub(/^[^0-9.]/,"",$2);print "https://" $2 ":6443"}' | tr '\n' ',')

mkdir -p /var/lib/kubelet/
echo "apiVersion: v1
kind: Config
clusters:
- cluster:
    certificate-authority: /var/lib/kubernetes/ca.pem
    server: ${CONTROLLER0}
  name: kubernetes
contexts:
- context:
    cluster: kubernetes
    user: kubelet
  name: kubelet
current-context: kubelet
users:
- name: kubelet
  user:
    token: $COMMON_TOKEN" > /var/lib/kubelet/kubeconfig
```

Create the kubelet systemd service file:

```
sh -c 'echo "[Unit]
Description=Kubernetes Kubelet
Documentation=https://github.com/GoogleCloudPlatform/kubernetes
After=docker.service
Requires=docker.service

[Service]
EnvironmentFile=/etc/dockerenv.conf
ExecStart=/usr/bin/kubelet \
  --allow-privileged=true \
  --api-servers=${API_SERVERS} \
  --cloud-provider= \
  --cluster-dns=10.32.0.10 \
  --cluster-domain=cluster.local \
  --container-runtime=docker \
  --network-plugin=kubenet \
  --kubeconfig=/var/lib/kubelet/kubeconfig \
  --serialize-image-pulls=false \
  --tls-cert-file=/var/lib/kubernetes/kubernetes.pem \
  --tls-private-key-file=/var/lib/kubernetes/kubernetes-key.pem \
  --v=2

Restart=on-failure
RestartSec=5

[Install]
WantedBy=multi-user.target" > /etc/systemd/system/kubelet.service'
```

Enable and start kubelet

```
systemctl daemon-reload
systemctl enable kubelet
systemctl start kubelet
```

Verification

```
systemctl status kubelet --no-pager
```

kube-proxy

Create the kube-proxy service file:

```
echo "[Unit]
Description=Kubernetes Kube Proxy
Documentation=https://github.com/GoogleCloudPlatform/kubernetes

[Service]
ExecStart=/usr/bin/kube-proxy \
  --master=${CONTROLLER0} \
  --kubeconfig=/var/lib/kubelet/kubeconfig \
  --proxy-mode=iptables \
  --v=2

Restart=on-failure
RestartSec=5

[Install]
WantedBy=multi-user.target" > /etc/systemd/system/kube-proxy.service
```

Enable and start kube-proxy:

```
systemctl daemon-reload
systemctl enable kube-proxy
systemctl start kube-proxy
```

Verification

```
systemctl status kube-proxy --no-pager
```

NOTE | Remember to run these steps on all workers.

Configuring the Kubernetes Client - Remote Access

Download and Install kubectl

OS X

```
wget -P /usr/local/bin/ https://storage.googleapis.com/kubernetes-
release/release/v1.5.2/bin/darwin/amd64/kubectl
chmod +x /usr/local/bin/kubectl
```

```
wget -P /usr/local/bin/ https://storage.googleapis.com/kubernetes-
release/release/v1.5.2/bin/linux/amd64/kubectl
chmod +x /usr/local/bin/kubectl
```

Configure Kubectl

Configure the kubectl client to point to the [Load Balancer](#) (if configured), or do this from the bastion and configure it to point to a controller. Set KUBERNETES_PUBLIC_ADDRESS to the IP address of the desired endpoint.

Recall the [common token](#) you set up.

Also be sure to locate the CA certificate from [Setting up a CA and TLS Cert Generation](#). Since we are using self-signed TLS certs we need to trust the CA certificate so we can verify the remote API Servers.

Build up the kubeconfig entry:

```
kubectl config set-cluster kubernetes-the-hard-way \
  --certificate-authority=ca.pem \
  --embed-certs=true \
  --server=https://${KUBERNETES_PUBLIC_ADDRESS}:6443

kubectl config set-credentials admin --token ${COMMON_TOKEN}

kubectl config set-context default-context \
  --cluster=kubernetes-the-hard-way \
  --user=admin

kubectl config use-context default-context
```

At this point you should be able to connect securly to the remote API server.

```
kubectl get componentstatuses
```

NAME	STATUS	MESSAGE	ERROR
controller-manager	Healthy	ok	
scheduler	Healthy	ok	
etcd-2	Healthy	{"health": "true"}	
etcd-0	Healthy	{"health": "true"}	
etcd-1	Healthy	{"health": "true"}	

```
kubectl get nodes
```

NAME	STATUS	AGE
fb011ce0-be49-4f5d-8d13-86153cdf42f7	Ready	7m
6bc8cb95-c927-4346-8ec8-e4e821dc4b22	Ready	5m
e8e8402e-3326-4d66-b3db-d62e588ac347	Ready	2m

Managing the Container Network Routes and Overlay Network

Now that each worker node is online, add an overlay network and routes so pods running on different machines can communicate.

These instructions require `kubect1` on the workers. Please follow [these instructions](#) to install `kubect1` on each worker.

Container Subnets

The IP addresses for each pod is allocated from the `podCIDR` range assigned to each Kubernetes worker through the node registration process.

The `podCIDR` is allocated from the cluster cidr range as configured on the Kubernetes Controller Manager with the following flag:

```
--cluster-cidr=10.200.0.0/16
```

Based on the above configuration each node receives a `/24` subnet. For example:

```
10.200.0.0/24
10.200.1.0/24
10.200.2.0/24
...
```

Populate the Routing Table

Populate the routing table with the L3 routes over the overlay network.

Use `kubect1` to print the `InternalIP` and `podCIDR` for each worker node.

```
kubect1 get nodes \
  --output=jsonpath='{range
  .items[*]}{.status.addresses[?(@.type=="InternalIP")].address} {.spec.podCIDR}
  {"\n"}{end}'
```

Output:


```
10.240.0.20 10.200.0.0/24
10.240.0.21 10.200.1.0/24
10.240.0.22 10.200.2.0/24
```

Create an overlay network

Fabric Networks do not forward multicasting so you need to create a unicast vxlan overlay. Create it so that it logically matches the podCIDR network.

NOTE Do this on each worker:

```
INTERNAL_IP=$(ip addr show net0 | awk '/inet /{gsub(/\./[0-9][0-9]/,"");print $2}')
VXLAN_IP=$(kubectl get nodes \
  --output=jsonpath='{range .items[*]}
  {.status.addresses[?(@.type=="InternalIP")].address} {.spec.podCIDR} {"\n"}{end}' |
  awk "/$INTERNAL_IP/{print \$2}" |
  sed -e 's/10\..200\../172.16./' -e 's@0/@1/@')
ip link add vxlan0 type vxlan id 1 dstport 0
eval $(kubectl get nodes \
  --output=jsonpath='{range
  .items[*]}{.status.addresses[?(@.type=="InternalIP")].address}{"\n"}{end}' |
  grep -v $INTERNAL_IP |
  awk '{ print "bridge fdb append to 00:00:00:00:00:00 dst " $1 " via net0" }')
ip addr add $VXLAN_IP dev vxlan0
ip link set up vxlan0
```

Create Routes

NOTE Do this on each worker:

```
eval $(kubectl get nodes \
  --output=jsonpath='{range
  .items[*]}{.status.addresses[?(@.type=="InternalIP")].address} {.spec.podCIDR}
  {"\n"}{end}' |
  grep -v $INTERNAL_IP |
  awk '{GW=$2;gsub("10.200","172.16",GW); gsub("0/24","1",GW); print "ip route add "
  $2 " via " GW }')
```

Deploying the Cluster DNS Add-on

Deploy the DNS add-on which is required for every Kubernetes cluster. Without the DNS add-on the following things will not work:

- DNS based service discovery

- DNS lookups from containers running in pods

Cluster DNS Add-on

Create the **kubedns** service:

```
kubectl create -f https://raw.githubusercontent.com/kelseyhightower/kubernetes-the-hard-way/master/services/kubedns.yaml
```

Verification

```
kubectl --namespace=kube-system get svc
```

NAME	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kube-dns	10.32.0.10	<none>	53/UDP,53/TCP	5s

Create the **kubedns** deployment:

```
kubectl create -f https://raw.githubusercontent.com/kelseyhightower/kubernetes-the-hard-way/master/deployments/kubedns.yaml
```

Verification

```
kubectl --namespace=kube-system get pods
```

NAME	READY	STATUS	RESTARTS	AGE
kube-dns-v19-965658604-c8g5d	3/3	Running	0	49s
kube-dns-v19-965658604-zwl3g	3/3	Running	0	49s

Smoke Test

Perform a quick smoke test to demonstrate that this cluster is working

Test

```
kubectl run nginx --image=nginx --port=80 --replicas=3
```

```
deployment "nginx" created
```

```
kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE
nginx-2032906785-ms8hw	1/1	Running	0	21s	10.200.2.2	worker2
nginx-2032906785-sokxz	1/1	Running	0	21s	10.200.1.2	worker1
nginx-2032906785-u8rzc	1/1	Running	0	21s	10.200.0.2	worker0

```
kubectl expose deployment nginx --type NodePort
```

```
service "nginx" exposed
```

*Grab the **NodePort** that was setup for the nginx service:*

```
NODE_PORT=$(kubectl get svc nginx --output=jsonpath='{range .spec.ports[0]}{.nodePort}')
```

*Grab the **EXTERNAL_IP** for one of the worker nodes:*

```
NODE_PUBLIC_IP=$(aws ec2 describe-instances \
  --filters "Name=tag:Name,Values=worker0" | \
  jq -j '.Reservations[].Instances[].PublicIpAddress')
```

*Test the nginx service using **cURL**:*

```
curl http://${NODE_PUBLIC_IP}:${NODE_PORT}
```

```
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
  }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
```

Congratulations. You now have a working cluster, built the hard way on Joyent.